

Introduction to Computer Vision in Python

Teodora Szasz

Image Analysis & Data Visualization Specialist

tszasz@uchicago.edu
[Research Computing Center](#)

What is Computer Vision?

Every picture tells a story



Every picture tells a story




Goal of computer vision is to write computer programs that can interpret images

Why study Computer Vision?

- **Images and movies are ubiquitous in both production and consumption**
- **Applications to manipulate images (movies) are a great need**
- **As are the systems to extract information from imagery**
 - **Surveillance**
 - **Building 3D models**
 - **Motion capture assisted**


What is the state of the art in
computer vision?





Workshop outline

Computational Models (Math)



Algorithm

Real Images

- **Introduction**
- **Filtering**
- **Template matching**
- **Edge detection**
- **Line detection - in a video**

What is an image?

Images as functions




Images as functions

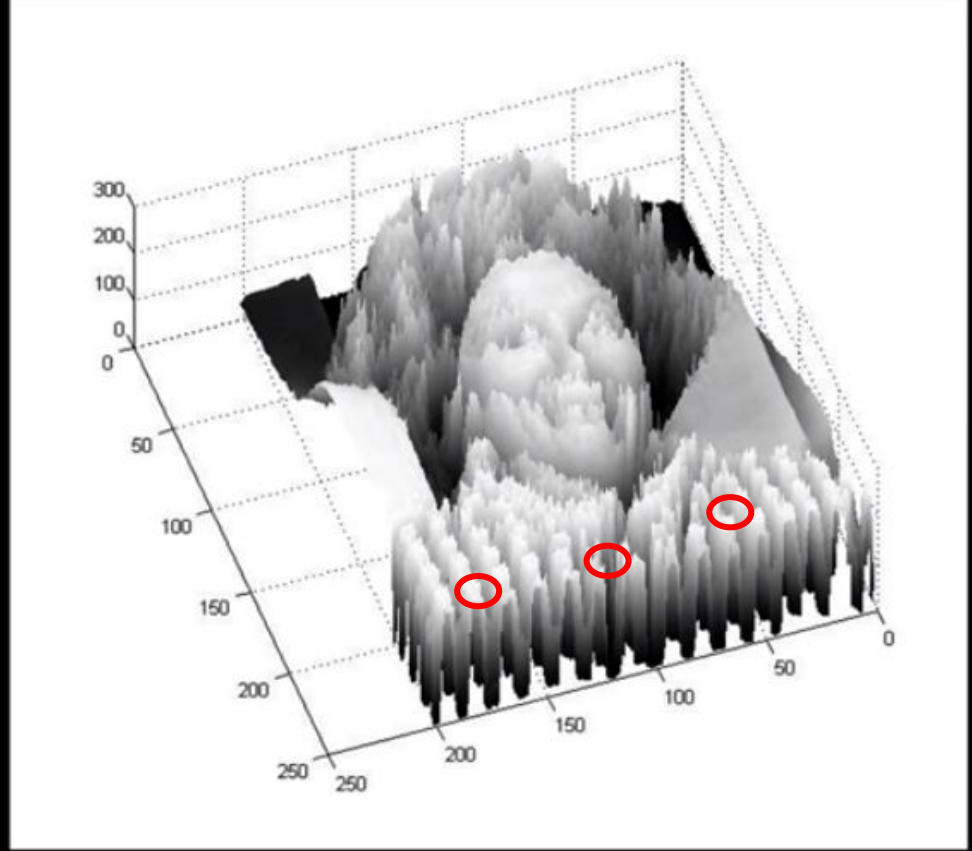


$I(x,y)$

Images as functions



Images as functions



Images as functions

$$f: [a,b] \times [c,d] \rightarrow [\min, \max]$$

Images as functions

$f: [a,b] \times [c,d] \rightarrow [\min, \max]$



0 255

Color images as functions

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

The real Phyllis




```
>> pd(40:60,30:40)
```

```
ans =
```

152	122	99	83	122	120	154	150	123	141	112
102	140	109	114	125	124	69	134	123	141	132
138	160	135	109	104	89	91	145	128	102	154
101	147	165	87	93	97	110	145	157	124	141
58	68	96	115	80	98	137	160	145	168	166
57	127	62	92	145	127	93	121	168	221	157
69	108	74	71	156	119	106	140	156	161	158
116	132	101	60	134	159	110	125	153	145	123
109	119	130	113	80	176	121	108	111	152	133
135	77	102	134	127	136	154	130	139	120	160
175	127	112	145	153	125	160	126	103	94	166
205	187	151	87	128	154	124	174	96	129	142
206	211	207	171	153	146	173	194	125	129	164
214	205	235	200	170	162	151	151	183	152	107
225	199	211	203	125	145	154	181	201	184	137
207	203	172	169	170	127	116	95	197	187	138
171	208	150	157	184	153	109	119	148	182	138
111	170	150	116	128	170	144	132	119	176	132
101	172	168	130	112	131	116	136	129	137	121
103	167	164	131	104	106	96	111	106	103	139
92	136	146	138	92	63	73	101	120	126	134

Hands-on: Gaussian Noise

1_Gaussian_Noise.ipynb




Hands-on: Gaussian Noise Image


2_Gaussian_Noise_Image.ipynb

How do we remove noise in an image?

Smoothing -1D




Weighted Moving Average




Weighted Moving Average in 2D

Reference
point

$F[x, y]$




$G[x, y]$





Weighted Moving Average in 2D

$$F[x, y]$$



$$G[x, y]$$



$$F[x, y]$$


Weighted Moving Average in 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Weighted Moving Average in 2D

...

$$F[x, y]$$
[illegible]

$$G[x, y]$$

[illegible]

Cross-correlation

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$H[u, v]$

$\frac{1}{9}$

1	1	1
1	?	1
1	1	1

“box filter”

$G[x, y]$

	0	10	20	30	30				

$$G = H \otimes F$$


Smoothing with a box filter



depicts box filter:
white = high value, black = low value



original



filtered

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$F[x, y]$

1	2	1
2	4	2
1	2	1



$H[u, v]$

This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



Smoothing with Gaussian

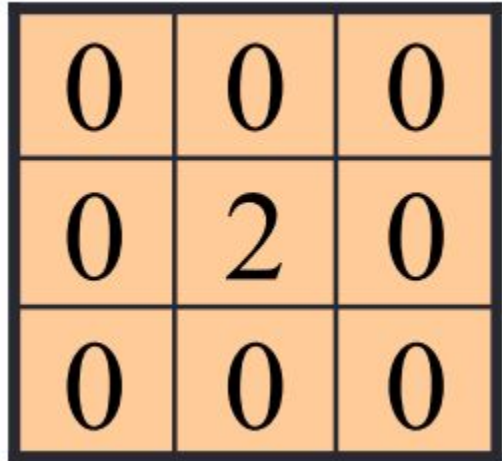


Hands-on: Gaussian Filter

3_Gaussian_Filter.ipynb

Hands-on: Linear Filter

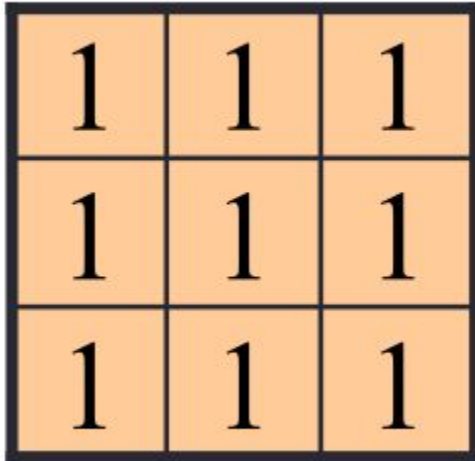
4_Linear_Filter.ipynb



0	0	0
0	2	0
0	0	0

—


$\frac{1}{9}$



1	1	1
1	1	1
1	1	1


Hands-on: Linear Filter

5_Median_Filter.ipynb




How filters will allow us to
extract features?

Template matching




Scene




Template (mask)

A toy example

Template matching



Detected template





Correlation map


Hands-on: Template Matching

6_Template_Matching.ipynb











Edges seem to be important ...






**Depth
discontinuity:
object boundary**

Discontinuity in surface orientation

**Reflectance
change:
appearance
information,
texture**

Cast shadow

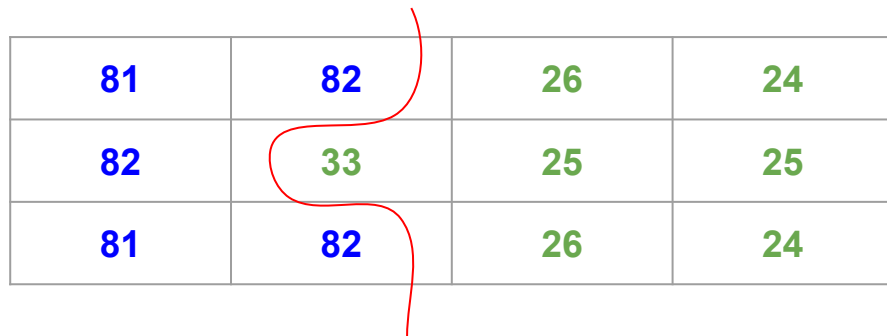
How we determine that a pixel at some location x,y is actually an edge pixel?



Edge detection

- **Idea**: look for a neighborhood with strong signs of change.

- **Questions**:
 - Neighborhood size
 - How to detect change




A 3x4 grid of numbers. The first column contains the value 81 in blue in all three rows. The second column contains 82 in blue in the top and bottom rows, and 33 in green in the middle row. The third and fourth columns contain the values 26, 25, and 24 in green, respectively, in all three rows. A red squiggly line highlights a neighborhood of change, starting from the top-right corner of the 82 cell, curving around the 33 cell, and ending at the bottom-right corner of the 82 cell.

81	82	26	24
82	33	25	25
81	82	26	24


Changes in a function? \Rightarrow Derivatives

- An edge is a place of rapid change in the image intensity function



Changes in a function? \Rightarrow Derivatives

- An edge is a place of rapid change in the image intensity function



How we will find the peaks?

Image

Differential
Operator

Image
gradient
function

Threshold
the **image**
gradient
function


Edge pixels

What is a gradient?


Image gradient

- The gradient of an image:


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$




$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Finite differences - discrete gradient

$$\frac{\partial f(x, y)}{\partial x} \approx f(x + 1, y) - f(x, y) \quad \text{"right derivative"}$$



**Finite difference
in x or in y?**



Finite differences - discrete gradient

$$\frac{\partial f(x, y)}{\partial x} \approx f(x + 1, y) - f(x, y) \quad \text{"right derivative"}$$



**Finite difference
in x or in y?**

Partial derivatives of an image



**In the images at bottom,
which one shows
changes in x direction
and which one in y
direction?**



Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---

(correlation filters)



$$\frac{\partial f(x, y)}{\partial y}$$

1
-1

Discrete gradient - operator

-1	1
----	---

0	0
-1	+1
0	0

0	0	0
-1/2	0	+1/2
0	0	0

-1	1
----	---

-1	1
----	---

-1	0	1
----	---	---

 / 2

average of “left” and
“right” derivative

Hands-on: Edge detection

7_Edge_Detection.ipynb

- Sobel operator: $\frac{1}{8} *$

-1	0	1
-2	0	2
-1	0	1

S_x


1	2	1
0	0	0
-1	-2	-1

S_y

Original image \longrightarrow Gradient magnitude \longrightarrow Thresholded gradient


- Other gradient masks: Sobel, Prewitt, Roberts

But, in the real world...



What we can apply to eliminate noise?

But, in the real world...



Apply Gaussian filtering => extract derivative of Gaussian filter


Hands-on: Canny edge operator

8_Canny_Edge_Operator.ipynb

1. **Filter image with the derivative of Gaussian**
2. **Find magnitude and orientation of the gradient**
3. **Non-maximum suppression:**
 - a. Thin multi-pixel wide “ridges” down to single pixel width
4. **Linking and thresholding(hysteresis):**
 - a. Define two thresholds: low and high
 - b. Use the high threshold to start edge curves and the low threshold to continue them

Hands-on: Canny edge operator

8_Canny_Edge_Operator.ipynb



Find the common edges in the two pictures

Image processing:



$$F : I(x,y) \rightarrow I'(x,y)$$

Computer vision:


$$F : I(x,y) \rightarrow \text{good stuff}$$

stuff e.g.: lines, circles, cars ==> called parametric models


Line fitting



Line fitting



Line fitting



How do we go from edges to lines?


Voting

- General technique where we let the features vote for all models that are compatible with it:
 1. Cycle through features, each casting votes for model parameters.
 2. Look for model parameters that receive a lot of votes.



Fitting lines

- **Hough Transform**

- Each edge point votes for compatible lines
- Look for lines that get many votes



Hough Transform



Hough Transform





image space



Hough (parameter) space






image space



Hough (parameter) space

Hough Transform

What line will be consistent with both points?



Hough Transform - the algorithm






image space



Hough (parameter) space

- Each edge point in image space votes for a set of possible parameters in Hough space


Hough Transform - the algorithm



- Each edge point in image space votes for a set of possible parameters in Hough space

Hough Transform - use a polar representation of lines

- Issues with usual (m,b) parameter space: can take on infinite values, undefined for vertical lines



d : perpendicular distance from line to origin

θ : angle the perpendicular makes with the x-axis

$$x \cos \theta - y \sin \theta = d$$

Hands-on: Hough transform

9_Hough_Transform.ipynb

Extensions:

- Using the gradient
 - $\theta = \text{gradient at } (x,y) \rightarrow \text{to reduce the voting time}$
- Can be used for detecting circles, squares, or any other shape

What is an image?

Until now: a function - a 2D pattern of intensity values

Now: a 2D projection of 3D points

- **Features and Matching**
 - **Corner detection**
 - **Panorama reconstruction**

<https://www.pyimagesearch.com/2016/01/11/opencv-panorama-stitching/>

Join the RCC's image analysis and
visualization [Slack team](#)